



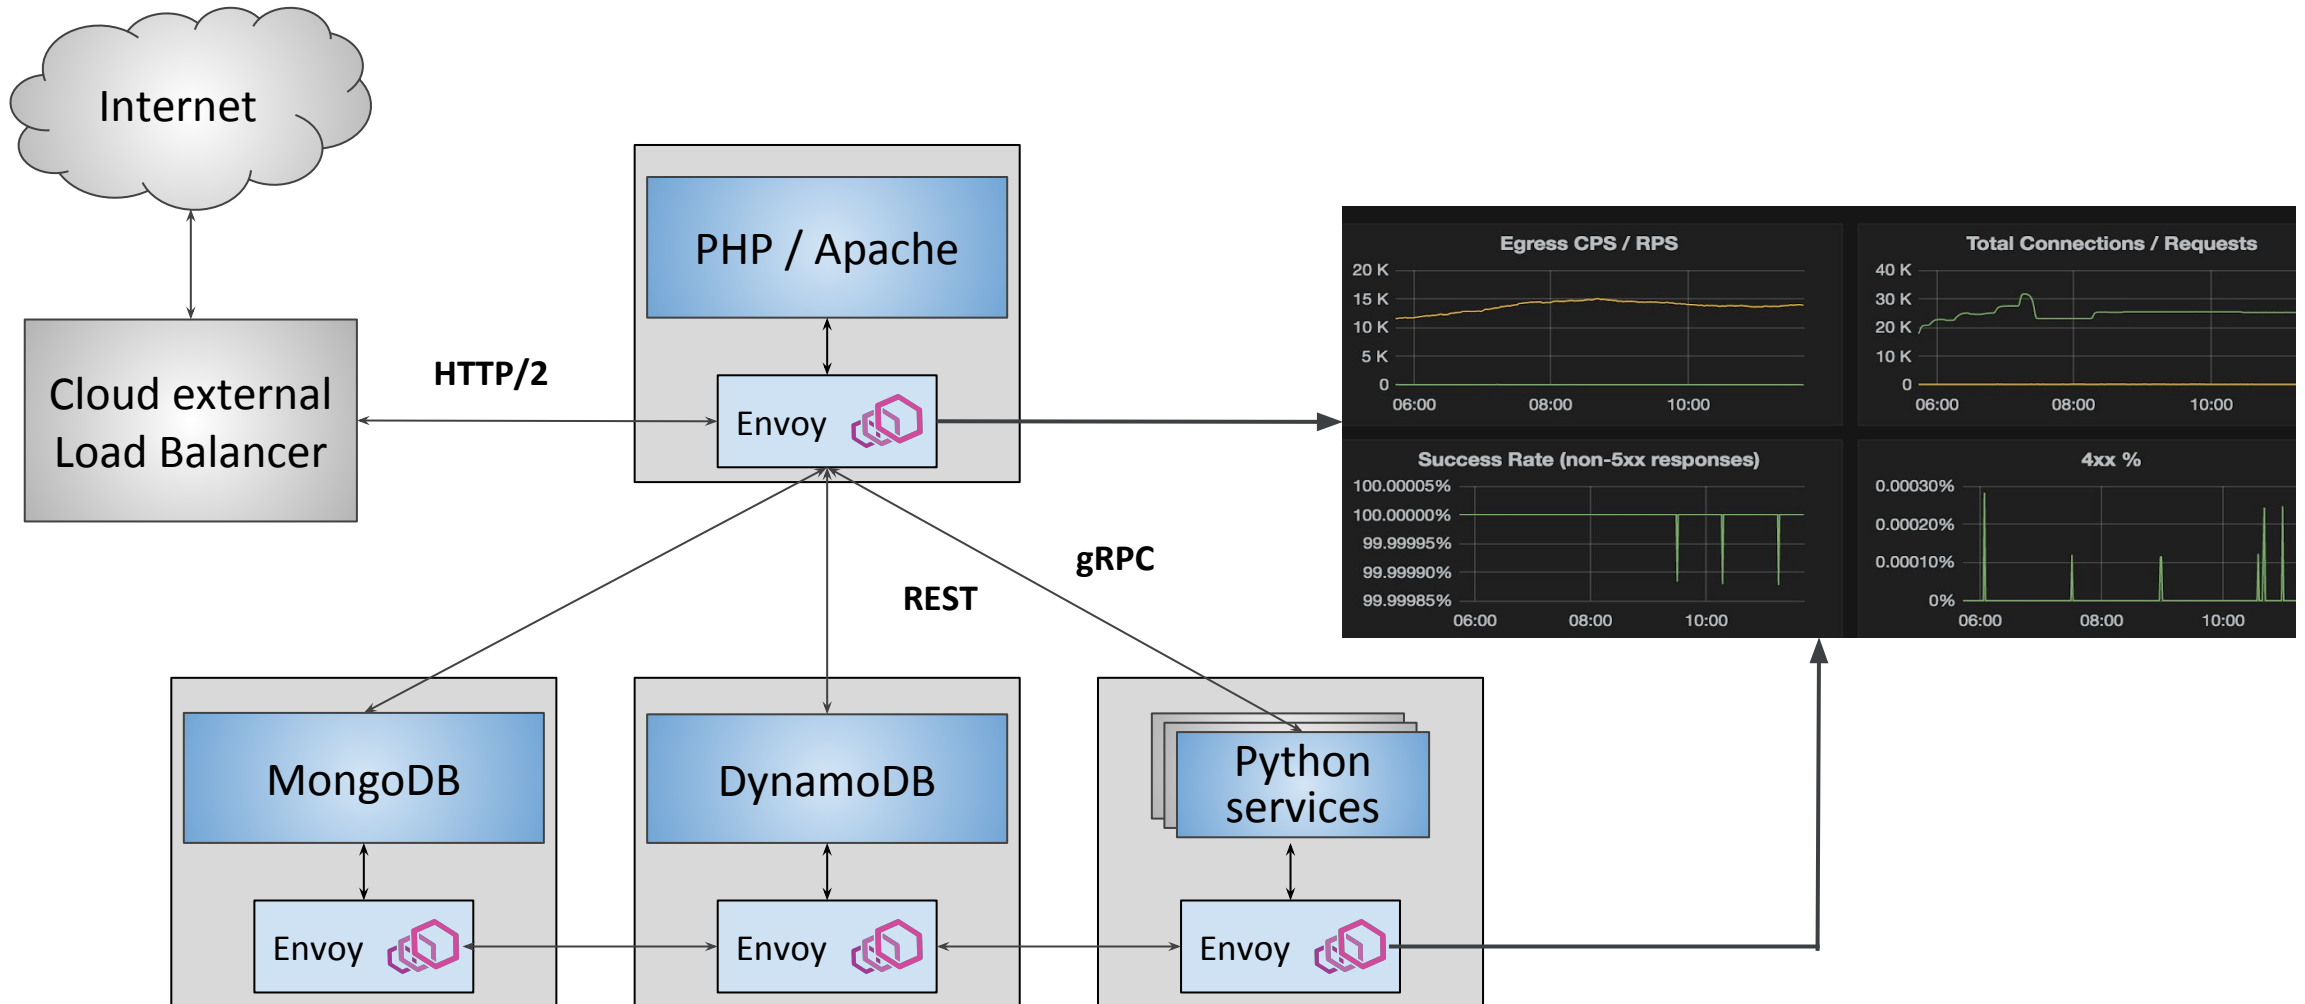
From Microbenchmarks to HTTP2 Load-testing: 5 Performance Tools and Techniques to Improve Envoy Scalability

Joshua Marantz (Google, [@jmarantz](#))

Otto van der Schaaf (We-Amp, [@oschaaf](#))

- **A Very Brief Intro to Envoy**
- Situational use of tools and techniques
 - Haphazard
 - Performance Annotations
 - Microbenchmarks
 - Application-level load testing
 - Fuzzing for Performance

Adding Observability to a Service Oriented Architecture



- **Best-in-class observability**
- **Performance**
- Reliability
- Modern codebase
- Extensibility
- Configuration API
- **Community**



So... how does
Envoy scale?



- Initialization performance
- Memory footprint
- Sensitivity to configuration size
- Sensitivity to traffic
- Sensitivity to large numbers of cores
- Data plane performance: Latency, Throughput

- Envoy as a front-end Proxy
- Cloud Providers
- CDNs

All these scenarios may define 10+k “clusters” (backend service groups), depending on scale

- **Haphazard**
- Performance Annotations
- Microbenchmarks
- Application-level load testing
- Fuzzing for Performance

- $O((100 * numClusters)^2)$ startup issue ([#2063](#))
 - Very easy to find using **haphazard** methods
 - ^C in the debugger
 - Or any other method of strobing active threads
- 20x speed up @10k clusters via better data structures ([#2358](#))
- >10 minutes → 22 seconds; still too slow
→ *mission not yet accomplished*

- Also easy-to-find haphazardly or via profiling: [#2373](#)
- Possible causes may include:
 - Regex package is slow (std::regex is much slower than RE2)
 - Too many regex lookups (lookup count found with [callgrind](#))
 - Very complicated regexes, some with [catastrophic backtracking](#)
- Regexes were used to add tagging structure to scalar stats
 - Needed for several stats sinks
- Solution was less obvious

Performance macro library, enabled by compile flag: [\[envoy proxy perf_annotation.h\]](#)

Initiates a performance operation, storing its state in perf_var.

PERF_OPERATION(perf_var)

Records performance data initiated with PERF_OPERATION.

PERF_RECORD(perf_var, category, description)

Dumps recorded performance data to stdout. Expands to nothing if not enabled.

PERF_DUMP()

```
for (extractor : all_tag_matcher_regexes_) {  
    PERF_OPERATION(perf);  
    std::smatch match;  
    if (std::regex_search(stat_name, match, extractor.regex) &&  
        match.size() > 1) {  
        /* Update tags, extract tags from stat_name */  
        PERF_RECORD(perf, "re-match", extractor.pattern);  
    } else {  
        PERF_RECORD(perf, "re-miss", extractor.pattern);  
    }  
}  
PERF_DUMP()
```

Annotated path;
static in code

Data-derived
annotation

10k clusters regex view

Initialization

3 regexes were much more expensive per/eval than the others

Almost all high-cost regex operations were mismatches

Duration(us)	# Calls	Mean(ns)	StdDev(ns)	Min(ns)	Max(ns)	Category	Description
1252735	910190	1376	483.682	242	47638	re-miss	envoy.grpc_bridge_method
1233659	910190	1355	453.037	234	28921	re-miss	cipher_suite
1216013	910190	1335	452	232	54768	re-miss	envoy.grpc_bridge_service
941602	1820243	517	279.096	238	27092	re-miss	envoy.http_conn_manager_prefix
749437	910190	823	361.951	261	98978	re-miss	envoy.response_code
695964	910170	764	393.6	238	43740	re-miss	envoy.respo
.....elided 12 similar rows, each costing > 0.6 seconds real time.....							
624921	910190	686	318.941	234	28680	re-miss	envoy.virtu
485891	910190	533	278.86	239	41152	re-miss	envoy.worker
474777	910190	521	272.132	227	36750	re-miss	envoy.mongo
474536	910190	521	264.689	233	23060	re-miss	envoy.rds_rou
473125	910184	519	263.062	236	24118	re-miss	envoy.li
464885	910190	510	270.477	228	28184	re-miss	envoy.virtual_host
426139	910000	468	276.855	407	26796	re-match	envoy.cluster_name
110	1	535	226	1086	re-miss	envoy.cluster_name	
65	1	711	283	1791	re-match	envoy.http_conn_manager_prefix	
15	1	737	490	1066	re-match	envoy.response_code_class	
3	1	418	401	1411	re-match	envoy.listener_address	

Only one regex match took material amounts of time: 0.4 sec

Total mismatch time: 17 seconds, or 76% of init time

Only one regex match took material amounts of time: 0.4 sec

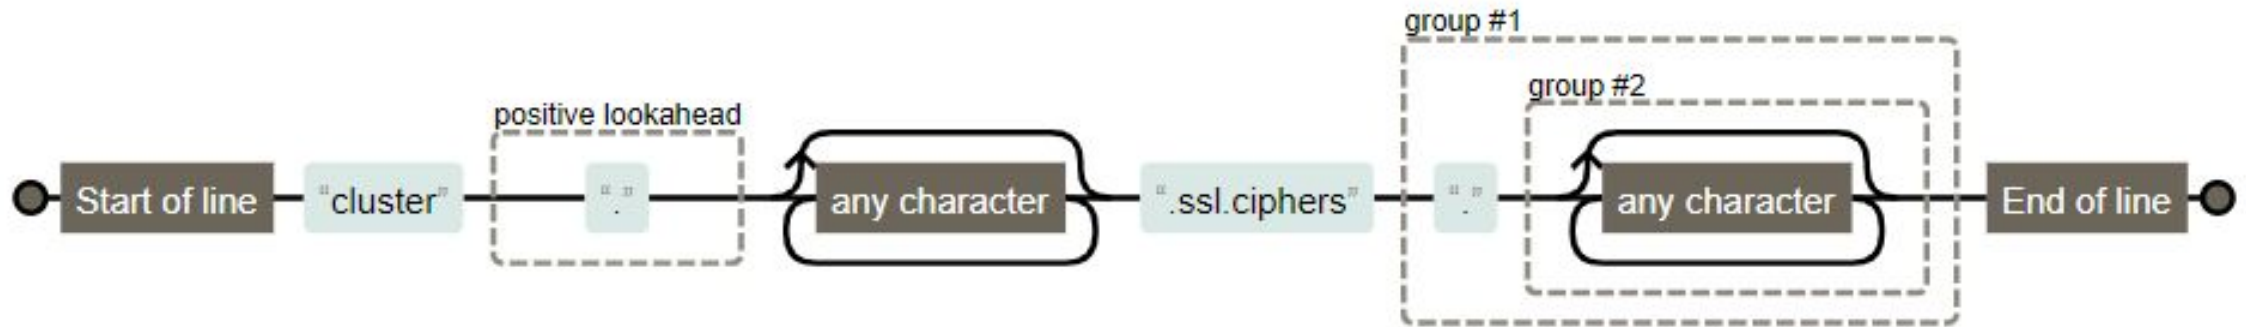
Total mismatch time: 17 seconds, or 76% of init time

Prefiltering to reduce the number of regex evals



Sample tag-extraction regex:

```
^cluster(=?\.).*?\.ssl\.ciphers(\.(*?))$
```



- Must begin with “**cluster.**” → build a map<prefix, RegexList>
- Must also contain “**ssl.ciphers**” → screen for substrings prior to regex

```
init() { computePrefixToRegexListMatch(); }
```

```
for (extractor : findPossibleExtractorsWithPrefix(stat_name)) {  
    PERF_OPERATION(perf);  
    if (extractor.substrMismatch(stat_name)) {  
        PERF_RECORD(perf, "re-skip", name_);  
        continue;  
    }  
}
```

```
std::smatch match;  
if (std::regex_search(stat_name, match, extractor.regex) &&  
    match.size() > 1) {  
    /* Update tags, tag_extracted name */  
    PERF_RECORD(perf, "re-match", extractor.pattern);  
} else {  
    PERF_RECORD(perf, "re-miss", extractor.pattern);  
    ...  
}
```

10k clusters with prefiltering: 3.75s initialization



Duration(us)	# Calls	Mean(ns)	StdDev(ns)	Min(ns)	Max(ns)	Category	Description
372381	910000	409	282.391	370	28669	re-match	envoy.cluster_name
122926	160054	768	335.647	487	44714	re-miss	envoy.response_code
121003	160034	756	422.152	482	28839	re-miss	envoy.response_code_class
41140	750136	54	82.8142	21	26367	re-skip	envoy.cipher
39580	750136	52	79.6177	21	15930	re-skip	envoy.cipher
36000	910000	39	77.9917	28	31885	re-skip	envoy.grpc_bridge_service
34498	910000	37	77.3431	27	25672	re-skip	envoy.grpc_bridge_method
32838	910000			27	16060	re-skip	envoy.grpc_bridge_method
66	137			33	5667	re-match	envoy.http_conn_manager_prefix
31	21			82	3788	re-miss	envoy.ssl_cipher
15	20	768	225.804	472	1470	re-match	envoy.response_code_class
11	15	747	169.845	614	1304	re-miss	envoy.listener_address
5	125	46	23.967	30	249	re-skip	envoy.dynamo_partition_id
5	125	44	16.2906	30	156	re-skip	envoy.dynamo_operation
5	125	42	11.1086	29	124	re-skip	envoy.rds_route_config
5	125	40	8.34395	28	65	re-skip	envoy.fault_downstream_cluster
4	125	39	10.2826	29	107	re-skip	envoy.dynamo_table
4	125	38	8.6811	30	112	re-skip	envoy.http_user_agent
3	6	643	418.244	415	1487	re-match	envoy.listener_address
0	9	70	11.9269	47	84	re-skip	envoy.http_conn_manager_prefix
0	9	33	15.1033	28	74	re-skip	envoy.worker_id

Most regexes
consume < 0.05
seconds

Most time-consuming
regex is now a MATCH

- Improving per-regex cost
 - Simplifying regular expressions so they evaluate faster
 - Switching to a better library, such as Google's RE2
 - `std::regex` allocates memory during pattern matching
 - Can run out of stack space on large input
- Enter Microbenchmarks ...



- A benchmark designed to measure the performance of a very small and specific piece of code.
 - [[google microbenchmark](#)]
 - [[alternative microbenchmark libraries](#)]
- Useful at multiple phases development
 - Weighing algorithm choices when writing code
 - Improving performance of a block of code known to be hurting performance

Microbenchmark for regexes



```
static void BM_StdRegex(benchmark::State& state)
{
    std::regex re("^cluster\\.((.*?)\\.)*");
    // alternate: ^cluster\\.([^\.\.]+)\.)*.*
    for (auto _ : state) {
        for (const std::string& cluster_input : inputs) {
            std::smatch match;
            if (std::regex_search(cluster_input, match, re)) {
                ++passes;
            }
        }
    }
}
```

BENCHMARK(BM_StdRegex);



https://github.com/jmarantz/envoy/blob/re-speed-benchmark/test/common/common/re_speed_test.cc

Microbenchmarks comparing variants



Benchmark	Time	CPU	Iterations
BM_StdRegex	1339 ns	1339 ns	441023
BM_StdRegexStringView	1288 ns	1288 ns	544246
BM_StdRegexStringViewAltPattern	1980 ns	1980 ns	354793
BM_RE2	1050 ns	1050 ns	666026
BM_RE2_AltPattern	448 ns	448 ns	1564090

- Test-patterns informed by given prefix-filtering (mismatches rare)
- `std::regex` improved 5% by using `string_view` rather than `std::string`
- RE2 on same regex improved another 25%
- RE2 with a better pattern provided a further 55% reduction, though that same pattern made `std::regex` worse

10% speed-up (3.75s -> 3.44s) with 1 new regex



Duration(us)	# Calls	Mean(ns)	StdDev(ns)	Min(ns)	Max(ns)	Category	Description
184379	910000	202	230.157	145	137020	re2-match	envoy.cluster_name
124345	160054	776	343.242	499	27504	re-miss	envoy.response_code
118753	160034	742	399.93	495	17136	re-miss	envoy.response_code
42086	750136	56	71.8637	24	16121	re-skip	envoy.response_code
39155	750136	52	76.7944	25	16226	re-skip	envoy.response_code
35589	910000	39	150.468	28	131877	re-skip	envoy.cipher_suite
33933	910000	37	65.1851	27	16145	re-skip	envoy.grpc_bridge_service
32949	910000	36	82.8183	27	25299	re-skip	envoy.grpc_bridge_method
61	137	451	399.466	244	3560	re-match	envoy.http_conn_manager_prefix
30	21	1469	494.582	1007	3295	re-miss	envoy.ssl_cipher
16	20	801	221.848	460	1427	re-match	envoy.response_code_class
11	15	757	137.208	637	1199	re-miss	envoy.listener_address
5	125	46	24.9203	30	242	re-skip	envoy.dynamo_partition_id
5	125	43	16.6237	30	155	re-skip	envoy.dynamo_operation
5	125	41	9.41567	29	101	re-skip	envoy.rds_route_config
5	125	40	15.0171	28	154	re-skip	envoy.fault_downstream_cluster
4	125	38	10.9989	28	112	re-skip	envoy.http_user_agent
4	125	38	10.9942	29	114	re-skip	envoy.dynamo_table
3	6	636	405.296	454	1463	re-match	envoy.listener_address
0	9	66	17.7795	30	92	re-skip	envoy.http_conn_manager_prefix

Pattern accelerated
from 0.37 seconds to
0.184 seconds

<https://github.com/envoyproxy/envoy/pull/8831>

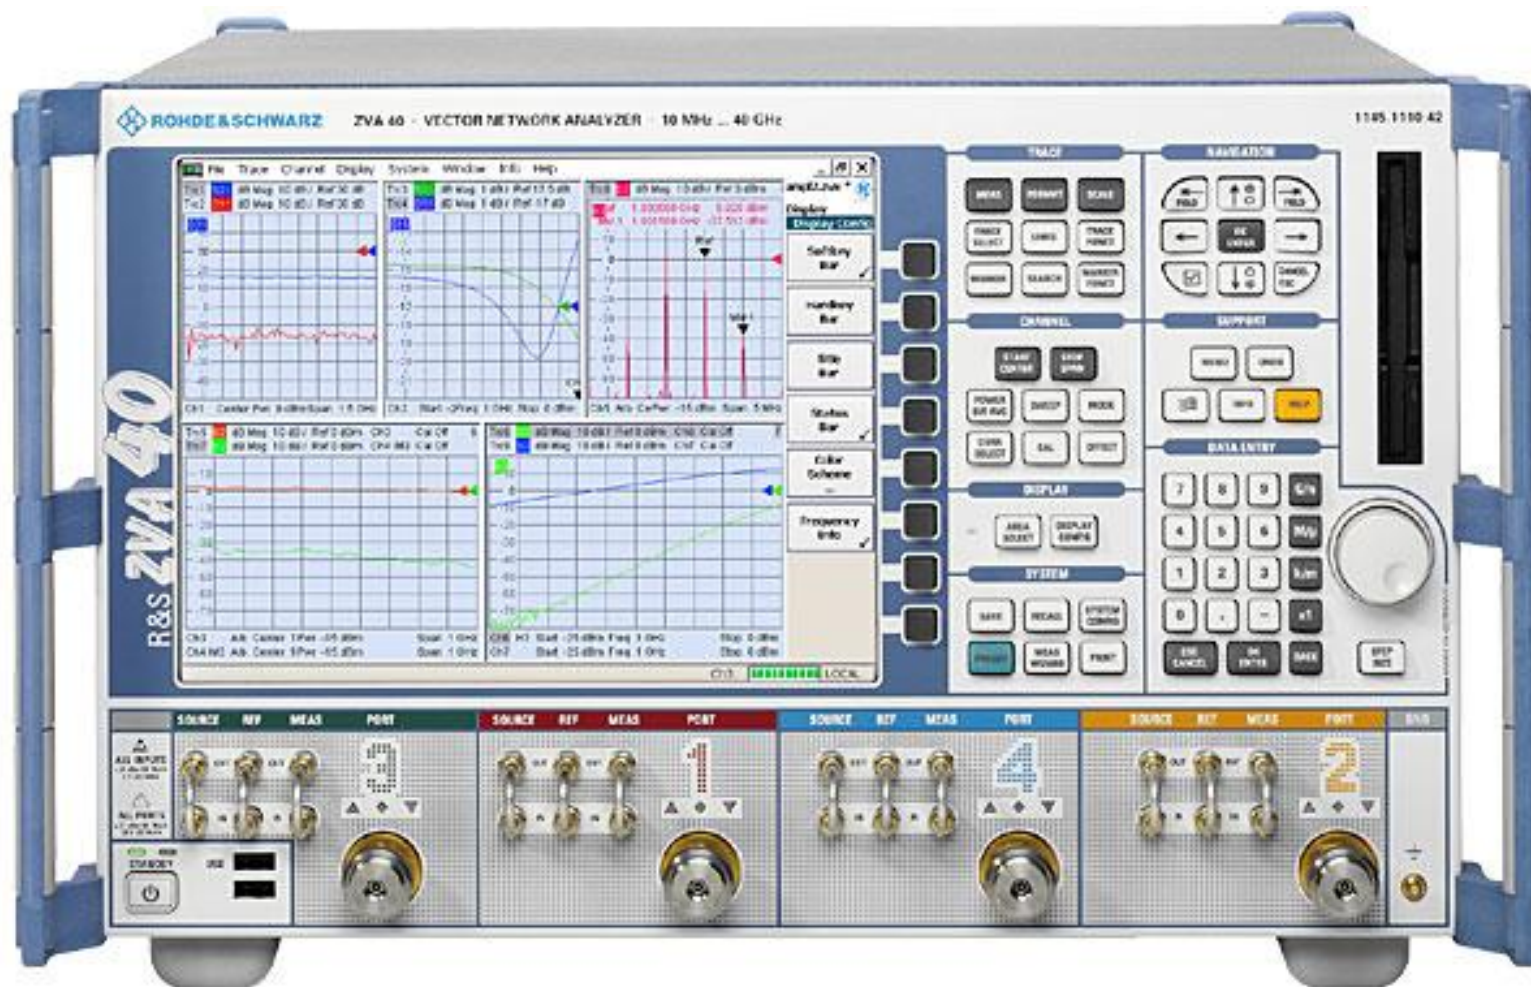
- ***Run the system to learn where the bottlenecks are***
- Annotate if needed to understand data-dependent issues
- Microbenchmark that data to converge to a solution
 - Easier to see small changes at nanosecond level
 - Faster iteration

- Haphazard
- Performance Annotations
- Microbenchmarks
- **Application-level load testing**
- Fuzzing for Performance

Performance-testing for an L7 Proxy

Testing tools must be \geq the performance of the system they are measuring:

- Capture long-tail latency
- 10-50 μ s resolution
- < 5% variance
- Native http2
- API/command-line



What HTTP(2)
load generators
are there?



- 10-50 μ s resolution
- Envoy-compatible engineering standards:
 - Test coverage, code reviews, continuous integration...
- HTTP/1 and HTTP/2 support, path to HTTP/3
- Open loop and closed loop modes
- Load-generation server via gRPC
- Latency histograms
- Structured output formatting
- OSS license compatible with Envoy's needs

Load Generation Evaluation funnel

Siege

Httpperf

Autobench

h2load

WRK

WRK2

JMeter

Web Polygraph

fortio

ab

Licensing

**Codebase & Cultural
Compatibility**

Accuracy

WRK2
fortio

- + Good at spreading load
- No HTTP2
- No test coverage
- Inactive community
- Sampling only
- Locking while sampling latency
- Millisecond granularity limitations

- + Supports HTTP2
- + Very nice visualizations
- Floating point math issues
- Go Garbage collection introduces jitter at μs granularity
- Governing scheduler (no fairness guarantees)
- Order of magnitude less accurate than WRK2 & Envoy Access logs control
- Uneven (batched) request-release timings

Decision: build a load-generator on top of Envoy




- + Supports HTTP2
- + Inherits HTTP3/Quic support when available in Envoy
- + High accuracy & scalability
- + Benefits from years of internal load-generation experience
- + Can establish great testing / code review / CI culture
- + Visualization agnostic
- Building takes time
- Building costs money
- Will require continued investment

Nighthawk: a
load generator
based on Envoy's
Network stack



- Built on Envoy network stack
- Github repo (parallel to Envoy, dependent on it)
- Envoy style test coverage, CI, C++ style, code reviews
- Performance knobs
 - Http2 vs Http
 - Max Active Requests
 - Concurrency
 - Targeted requests-per-second
- Outputs latency histograms, other stats

Nighthawk: closed-loop mode



- Resource limits will induce back-pressure and influence request timings
- Wait time for configured resource limits reported as “blocking”

Nighthawk: command line closed-loop test



```
taskset -c 0 ./nighthawk_client --concurrency auto --rps 15000 --duration 10 --connections 1 127.0.0.1:10000
```

```
[10:11:30.169193][5235][I] [source/client/process_impl.cc:170] Detected 1 (v)CPUs with affinity..
```

....

Request start to response end

samples: 149942

mean: 0s 000ms 037us

pstdev: 0s 000ms 003us

Percentile	Count	Latency
0	1	0s 000ms 034us
0.5	74976	0s 000ms 037us
0.75	112471	0s 000ms 037us
0.8	120045	0s 000ms 037us
0.9	134980	0s 000ms 038us
0.95	142447	0s 000ms 041us
0.990625	148541	0s 000ms 049us
0.999023	149796	0s 000ms 067us
1	149942	0s 000ms 572us

Blocking. Results are skewed when significant numbers are reported here.

samples: 241

mean: 0s 000ms 042us

pstdev: 0s 000ms 006us

Percentile	Count	Latency
0	1	0s 000ms 037us
0.5	121	0s 000ms 040us
0.75	181	0s 000ms 041us
0.8	193	0s 000ms 041us
0.9	217	0s 000ms 045us
0.95	229	0s 000ms 050us
0.990625	239	0s 000ms 054us
1	241	0s 000ms 128us

Counter	Value	Per second
benchmark.http_2xx	149943	14994.29
upstream_cx_http1_total	1	0.10
upstream_cx_rx_bytes_total	24440709	2444069.14
upstream_cx_total	1	0.10
upstream_cx_tx_bytes_total	8996580	899657.35
upstream_rq_total	149943	14994.29

Nighthawk: open-loop mode



- No feedback-loop based on resource limits. Requests are released unconditionally when they are due.
- Counters will track failures because of configured Nighthawk resource limits.

Nighthawk: command line open-loop test



```
taskset -c 0 ./nighthawk_client --concurrency auto --rps 15000 --duration 10 --connections 1 --open-loop 127.0.0.1:10000
```

```
[10:19:10.081859][6043][I] [source/client/process_impl.cc:170] Detected 1 (v)CPUs with affinity..
```

```
...
```

```
Request start to response end
```

```
  samples: 149849
```

```
  mean:    0s 000ms 038us
```

```
  pstdev:  0s 000ms 004us
```


Percentile	Count	Latency
0	1	0s 000ms 036us
0.5	75085	0s 000ms 038us
0.75	112526	0s 000ms 038us
0.8	120009	0s 000ms 038us
0.9	134915	0s 000ms 038us
0.95	142360	0s 000ms 039us
0.990625	148446	0s 000ms 043us
0.999023	149703	0s 000ms 061us
1	149849	0s 001ms 408us

Counter	Value	Per second
benchmark.http_2xx	149850	14984.97
upstream_cx_http1_total	1	0.10
upstream_cx_overflow	6	0.60
upstream_cx_rx_bytes_total	24425550	2442549.56
upstream_cx_tx_bytes_total	8991000	899098.00
upstream_rq_total	149850	14984.97

Nighthawk: zooming in

Export to **Fortio** for close inspection of the long tail. It's unexpectedly long.

```
./nighthawk_client .. --output-format fortio > export-to-fortio.json
```

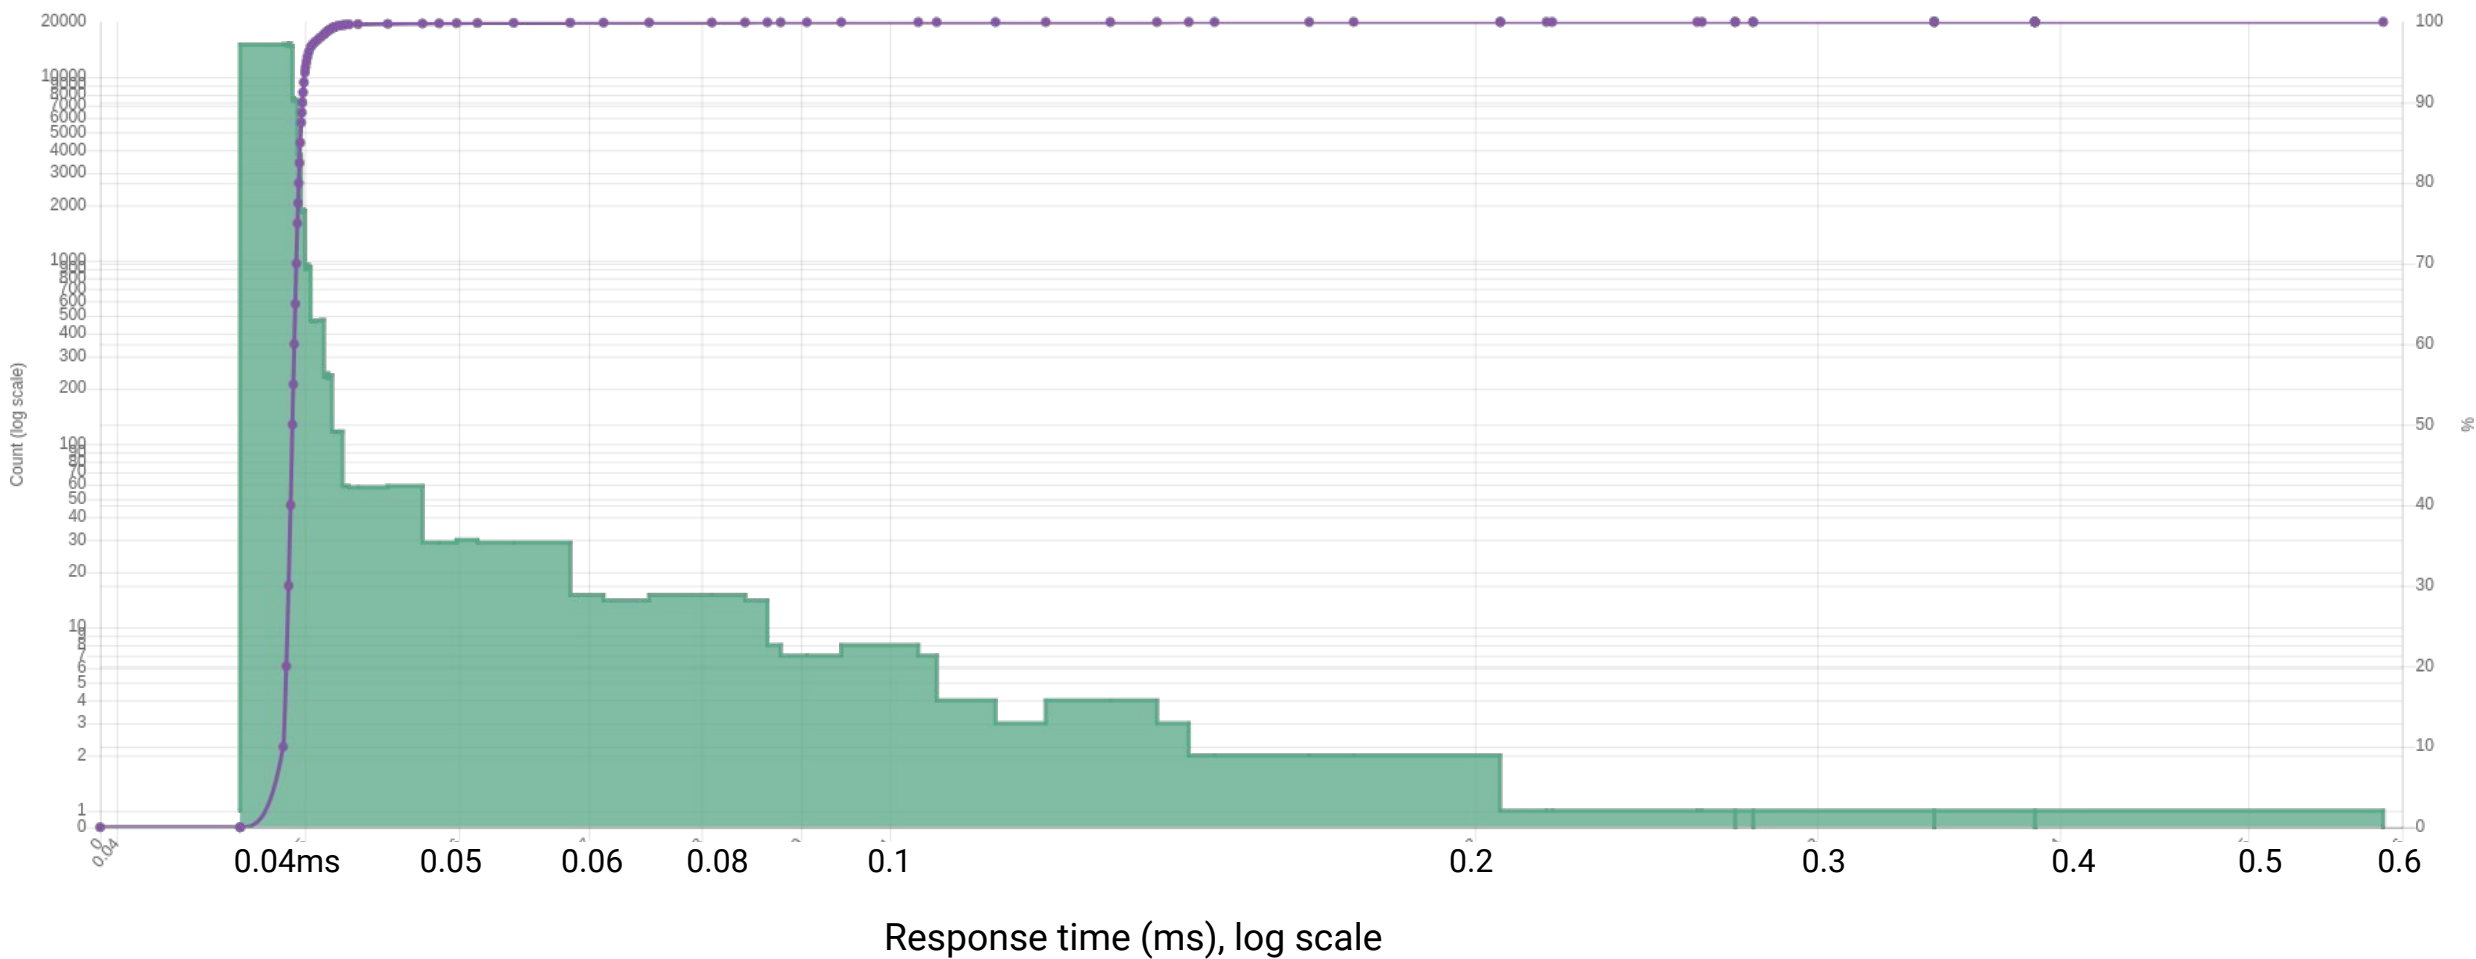


Nighthawk: single threaded pinned test server



Pinning test-server to a CPU has a dramatic impact

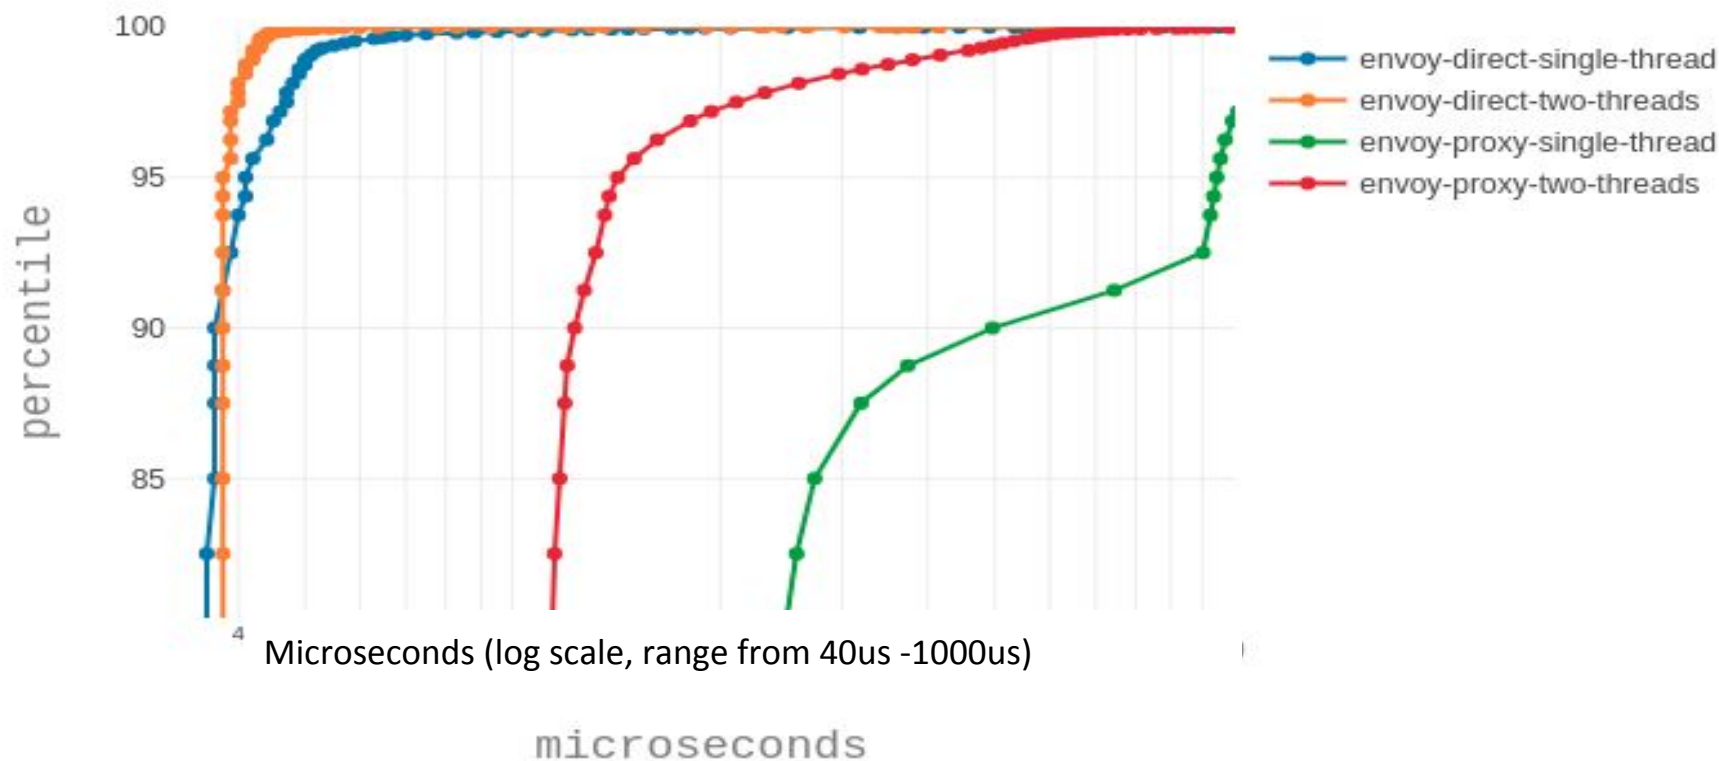
```
taskset -c 39 bazel-bin/nighthawk_test_server -c ~/envoy.yaml
```



Nighthawk: Measuring small differences

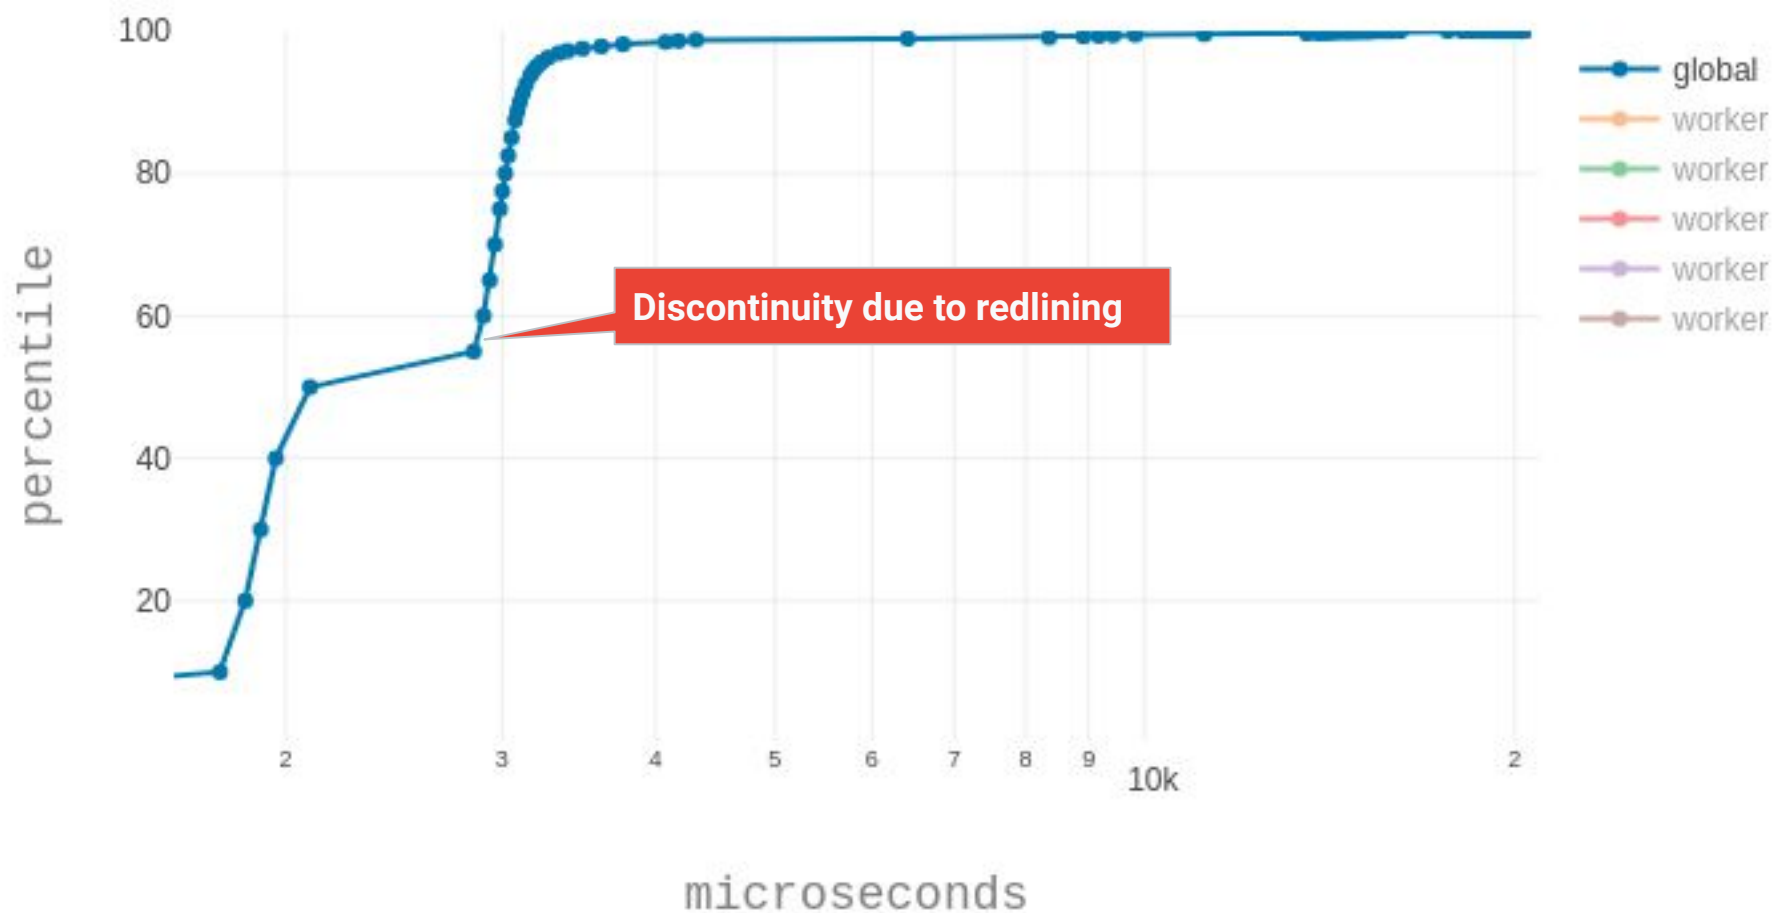


- Envoy in direct response mode vs. through Envoy in proxy mode
- Single-threaded Envoy vs double-threaded Envoy
- Small request / reply size, H1, perfect keep-alive @ 15K RPS



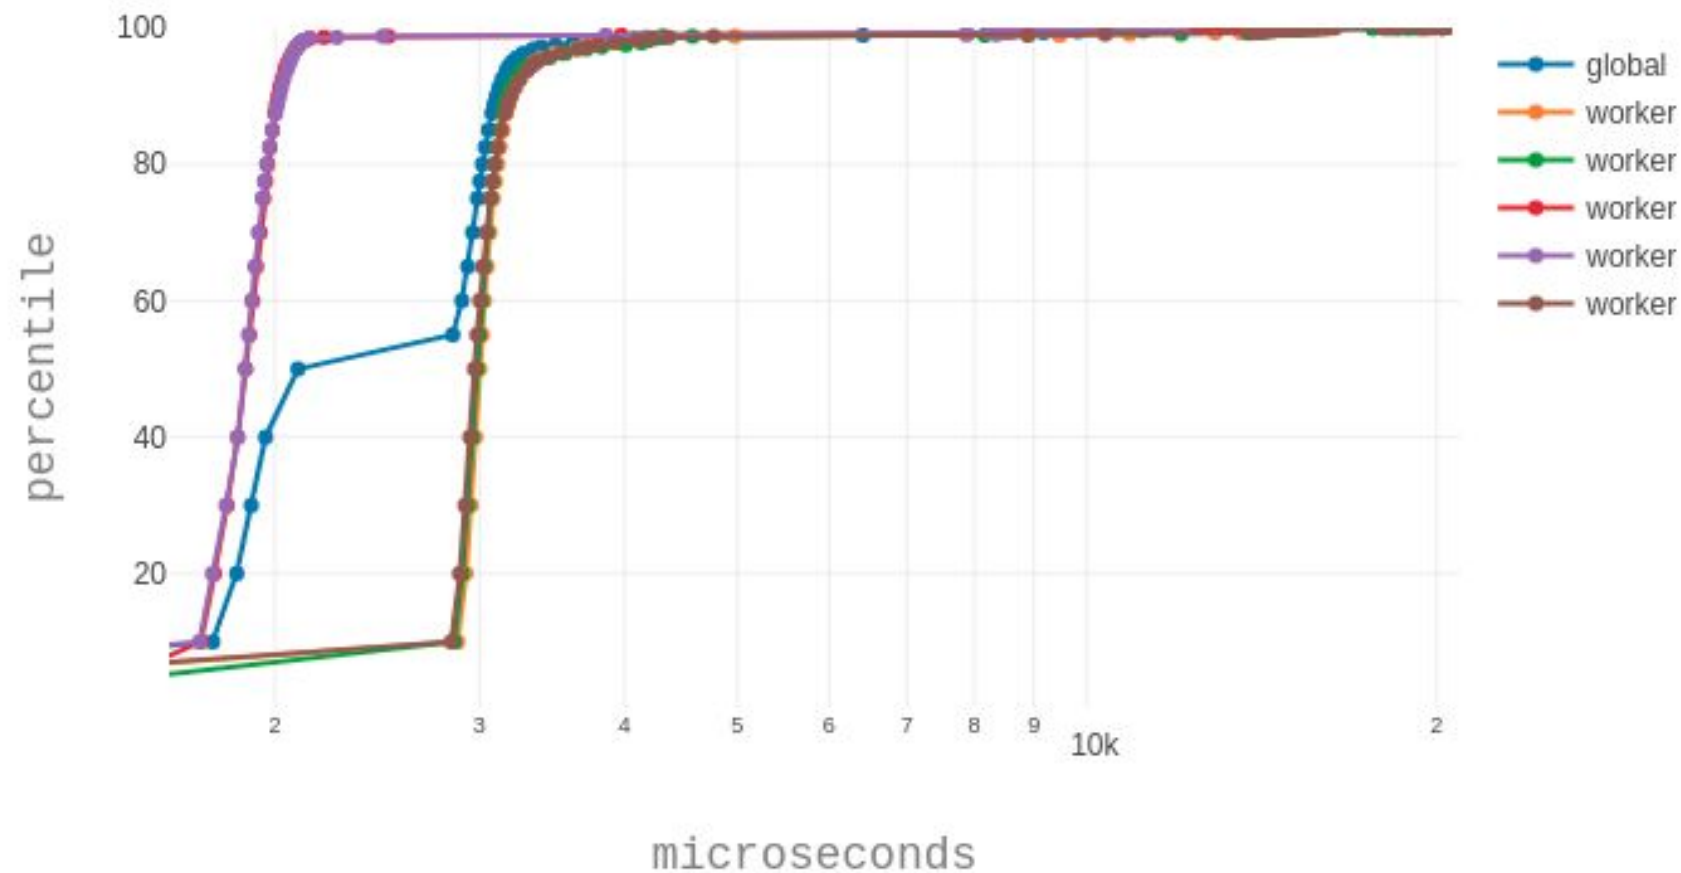
Nighthawk: per-worker reporting

- Redline testing example: server on 2 cpu cores
- Nighthawk on 5 cores



Nighthawk: per-worker reporting

The per worker visualization show imbalances, explaining the odd shape of the aggregated result.
Potential backend process hotspotting.



Nighthawk: Envoy H1 vs H2



Nighthawk's facilitates writing targeted benchmarks with a small python framework .
This will capture profiling data as well as yield structured output in json.

Let's compare downstream H1 vs H2(C) performance for small request / replies via closed-loop high rps tests.

Sample test scripts

```
def test_h1_concurrent_redline(http_test_server_fixture):
    http_test_server_fixture.test_server.enableCpuProfiler()
    parsed_json, _ = http_test_server_fixture.runNighthawkClient(
        [http_test_server_fixture.getTestServerRootUri(), "--rps", "50000", "--duration", "60", "--concurrency 120",
        "--connections", "100", "--max-pending-requests", "100"])
```



```
def test_h2_concurrent_redline(http_test_server_fixture):
    http_test_server_fixture.test_server.enableCpuProfiler()
    parsed_json, _ = http_test_server_fixture.runNighthawkClient(
        [http_test_server_fixture.getTestServerRootUri(), "--rps", "50000", "--duration", "60", "--concurrency 120", "--h2",
        "--max-pending-requests", "100"])
```

Command to execute:

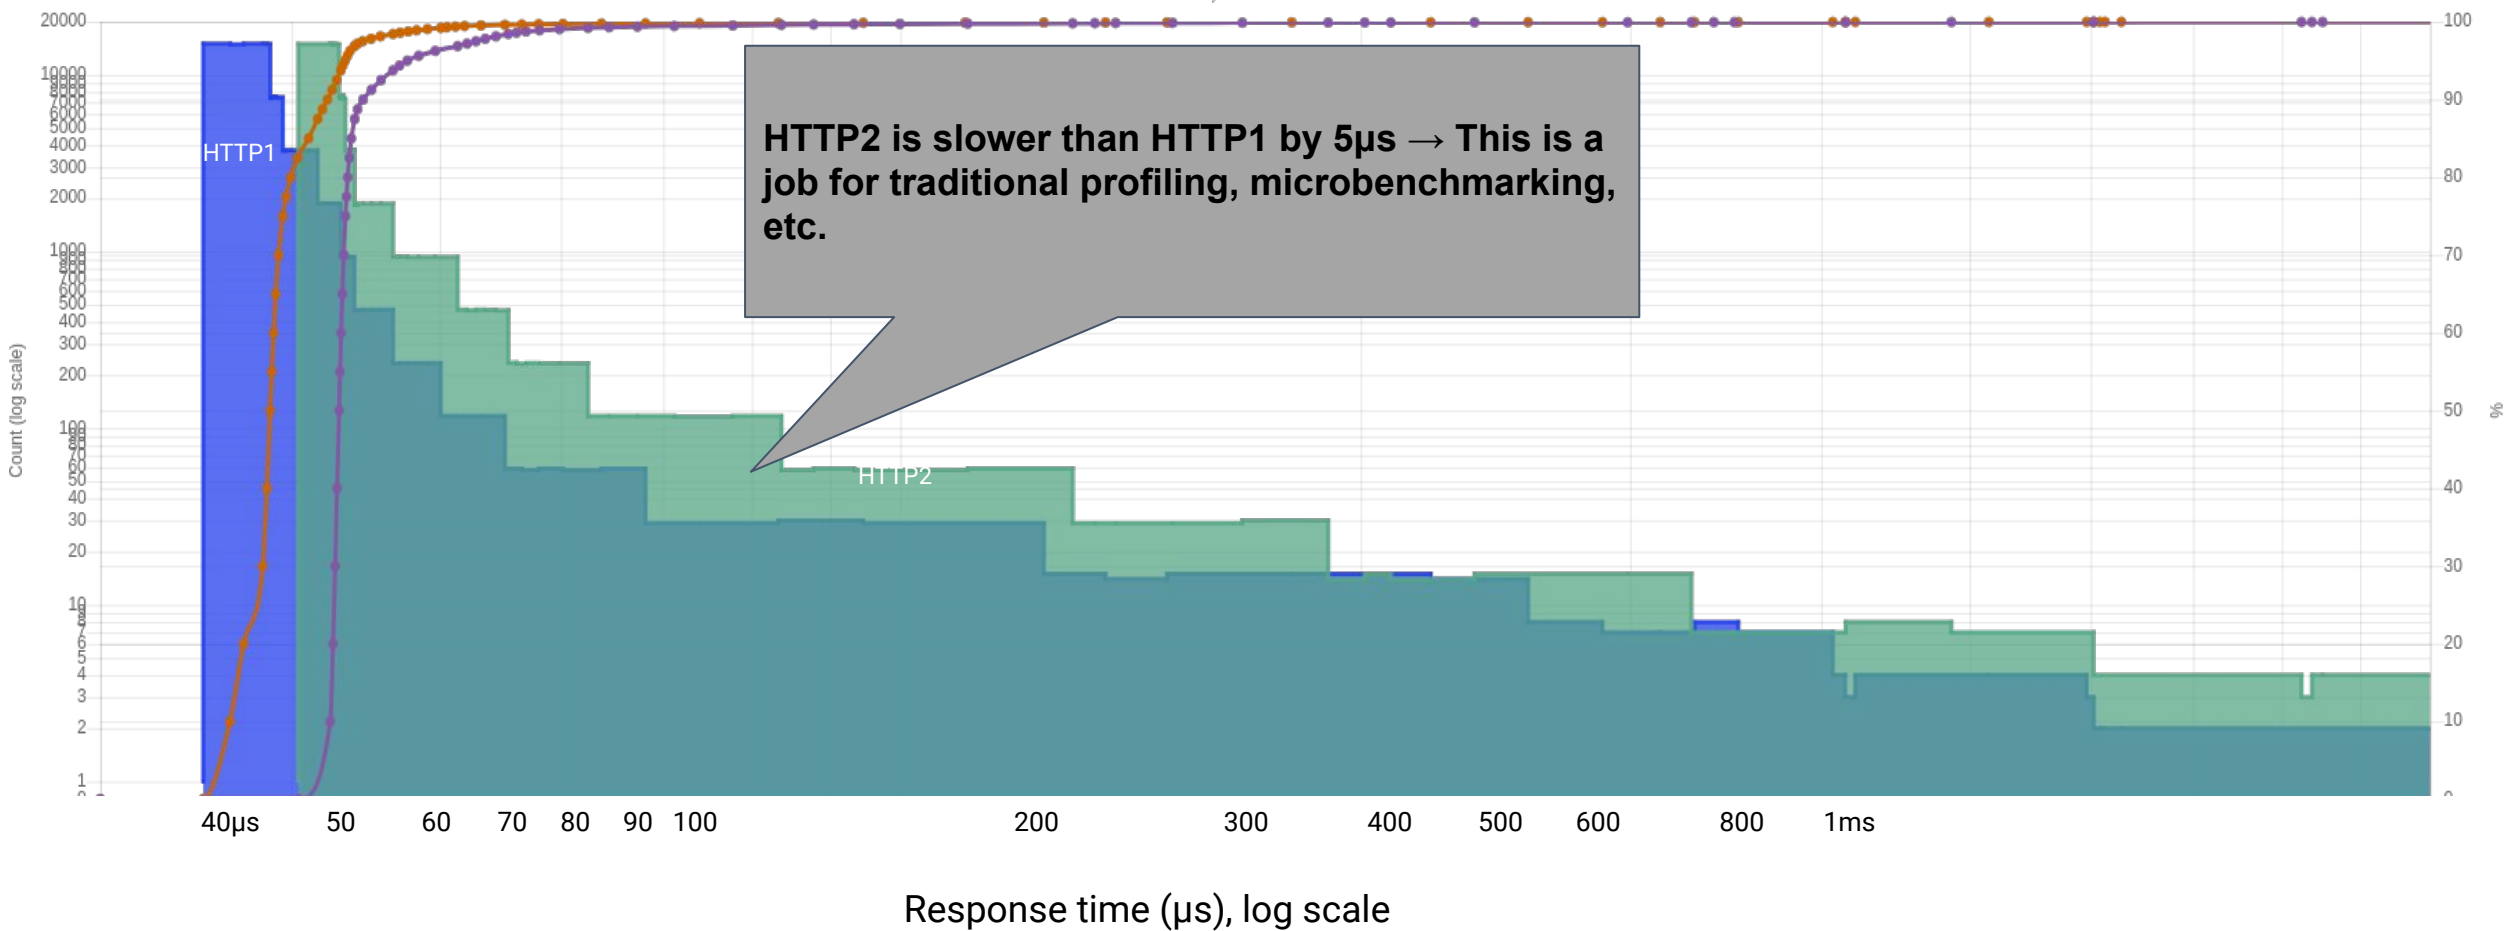
```
bazel test --test_env=ENVOY_IP_TEST_VERSIONS=v4only --test_env=HEAPPROFILE=
--test_env=HEAPCHECK= --cache_test_results=no --compilation_mode=opt --cxxopt=-g --cxxopt=-ggdb3
//benchmarks:*
```

HTTP1

```
nighthawk_client --rps 15000  
--connections 10 ...
```

HTTP2

```
nighthawk_client -h2 --rps 15000  
--max-active-requests 10 ...
```



Nighthawk: Envoy H1 vs H2 performance

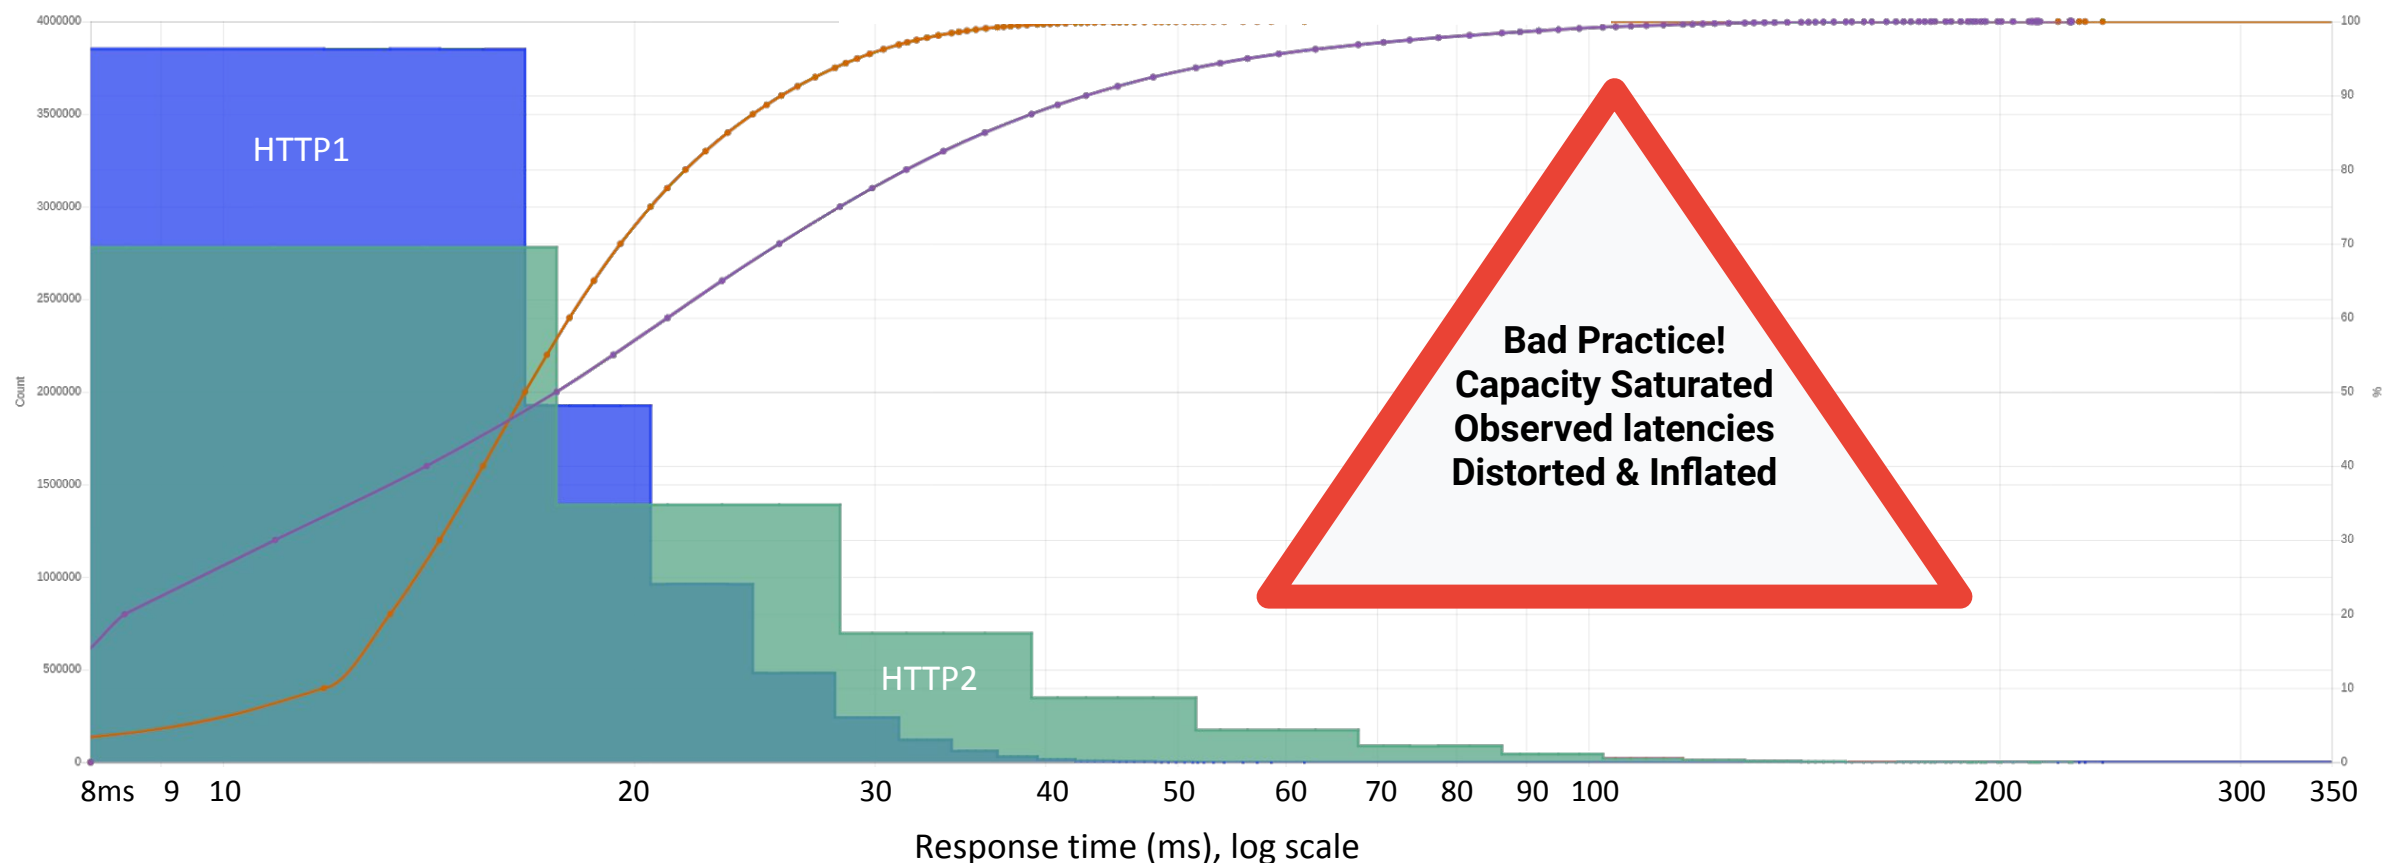


Bad practice: This test (more then) saturated the available computing capacity.

Observed latencies are distorted and inflated!

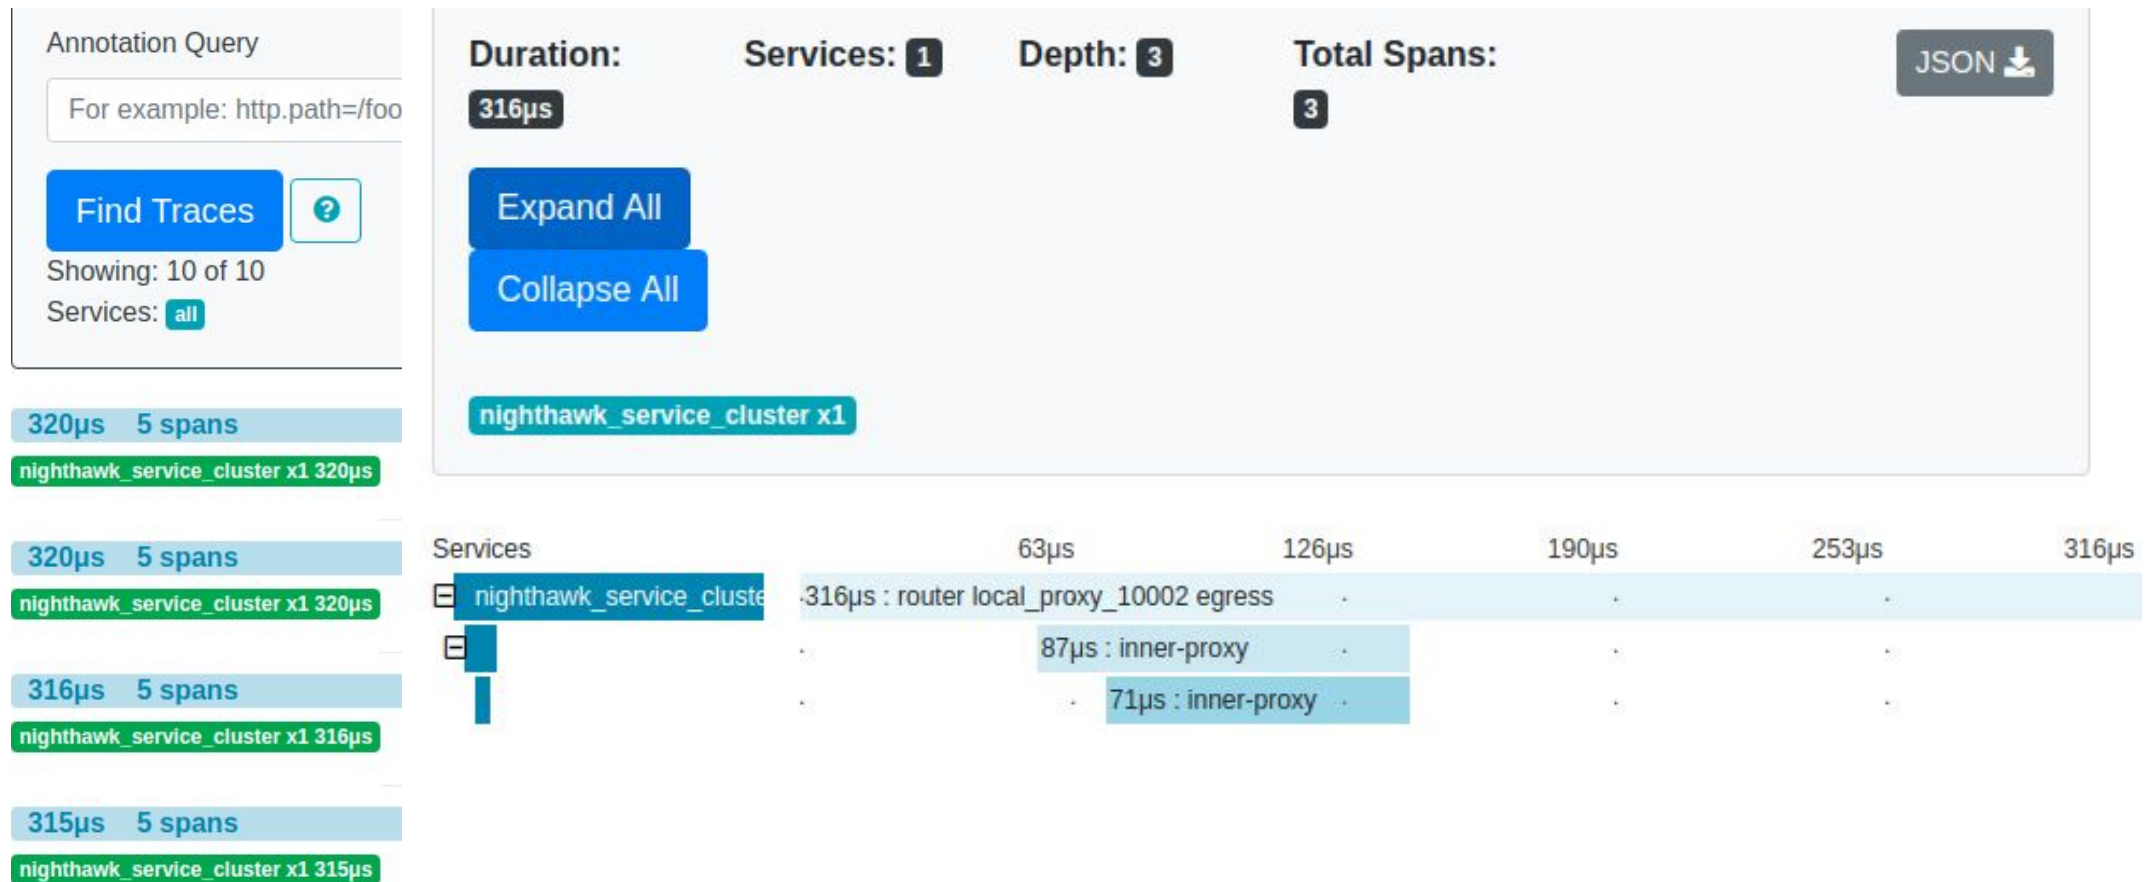
Max RPS, small request/reply - Envoy on 60 cores, 120 clients (Nighthawk threads) on 20 cores, single connection.

H1: ~641000 rps, H2 ~463000 rps



Nighthawk: tracing

Supports Zipkin trace initiation. Leverages Envoy's OpenTracing facilities.



- **Termination predicates:** specify when and how to terminate execution
- **Nighthawk as a service:** bi-directional streaming grpc service to request load tests and receive status updates
- **Replay:** pull to-be-replayed traffic from a grpc service (in review)

- Haphazard
- Performance Annotations
- Microbenchmarks
- Application-level load testing
- **Fuzzing for Performance**

- Fuzzers stimulate subsystems with random data
- Finds code-paths missed by unit, integration, and system tests
- Learns what patterns wake up new code and spends more effort varying those
- Security focus, but finds performance issues too

Finding Performance Problems from generated patterns



Fuzzer	Avg Exec/Second	Timeouts (%)	Regular Crash (%)
H1 capture direct response	1.9	3.8	7.3
H1 capture	2.8	1.2	77.7
Access Log Formatter	4.8	26	17.9
Conn Manager	4.8	0.8	22.4
New Buffer	9.3	0	0
Buffer	12.5	0	0

Key Takeaway: Performance data can be hard to find, but low executions/sec, high timeouts may be a signal

https://oss-fuzz.com/fuzzer-stats?group_by=by-fuzzer&date_start=2019-10-01&date_end=2019-10-04&fuzzer=libFuzzer&job=libfuzzer_asan_envoy&project=envoy


Finding Performance Problems from generated patterns



Issue 16325: envoy:h1_capture_direct_response_fuzz_test: Timeout in envoy_h1_capture_direct_

Reported by [ClusterFuzz-External](#) on Fri, Aug 9, 2019, 8:21 PM EDT

Project Member

 Only users with Commit permission can view this issue.

Detailed report: <https://oss-fuzz.com/testcase?key=5672448908853248>

Project: envoy

Fuzzing engine: libFuzzer

Fuzz target: h1_capture_direct_response_fuzz_test

Job Type: libfuzzer_ubsan_envoy

Platform Id: linux

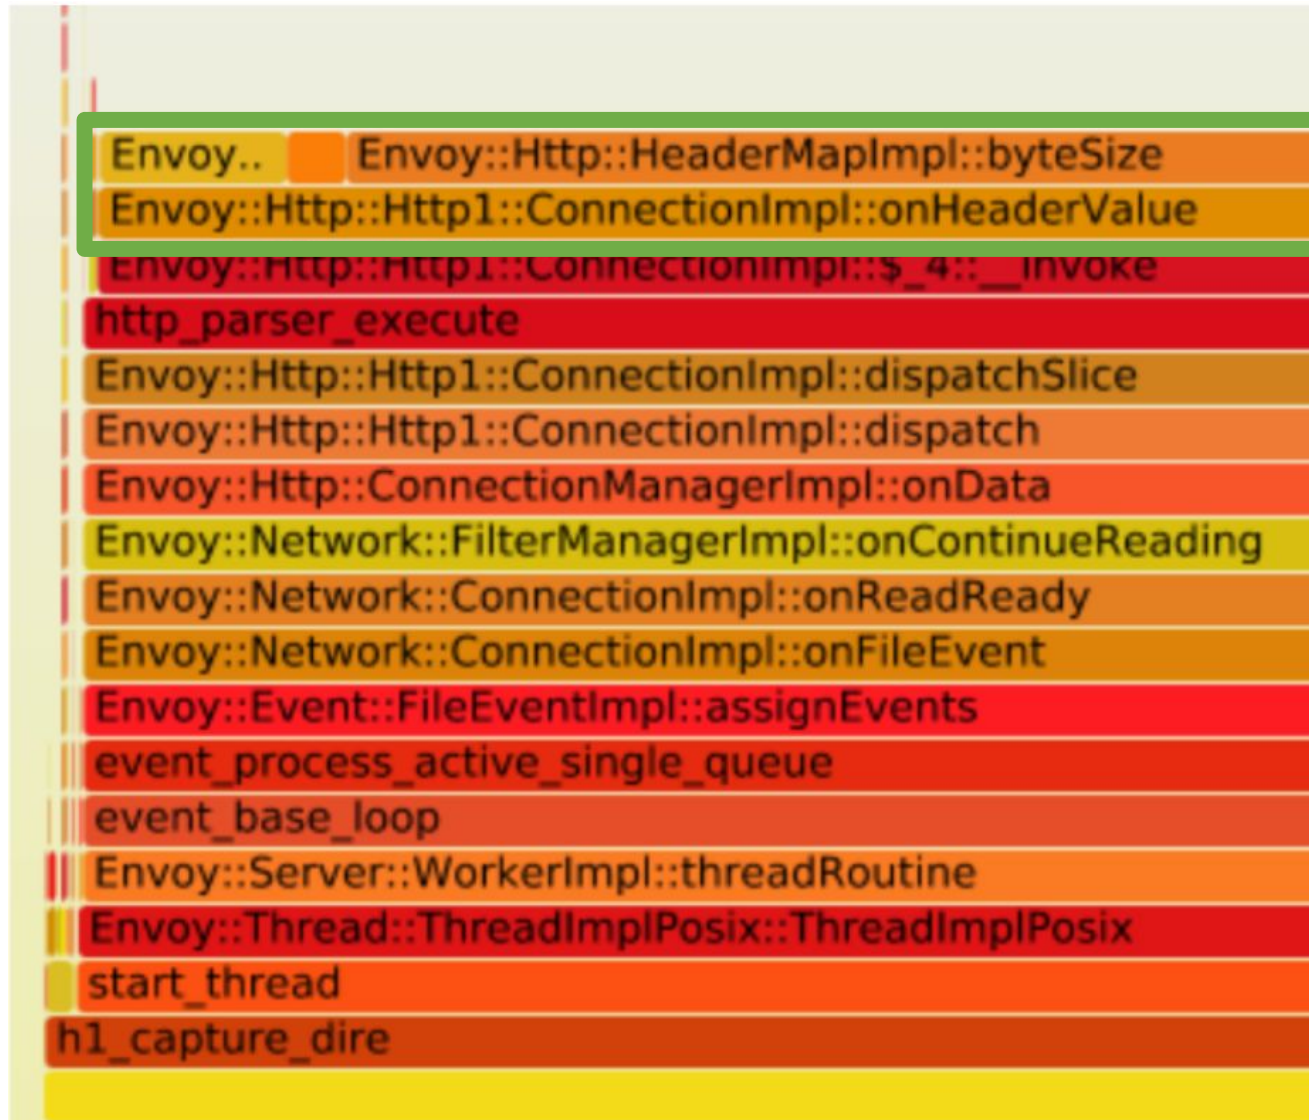
Crash Type: Timeout (exceeds 25 secs)

Crash Address:

Crash State:

envoy_h1_capture_direct_response_fuzz_test

Flame-Graph from repro of fuzz timeout



$O(n^2)$ update of
byte-size as new
headers arrive.


...and the resultant CVE

CVE-2019-15226

<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2019-15226>

Upon receiving each incoming request header data, Envoy will iterate over existing request headers to verify that the total size of the headers stays below a maximum limit. The implementation in versions 1.10.0 through 1.11.1 for HTTP/1.x traffic and all versions of Envoy for HTTP/2 traffic had **$O(n^2)$ performance** characteristics. **A remote attacker may craft a request** that stays below the maximum request header size but consists of many thousands of small headers to **consume CPU and result in a denial-of-service attack.**

Envoy Polling loop Latency vs. Percentile



- Most Envoy fuzz tests are very slow, masking problems in production code
- Assertion failures make fuzzing slow or crashy
 - Fuzzing will find more if the fuzzing tests avoid these
- Assertions themselves may be slow
- The tooling is not tuned for this use case at all, but there is opportunity for improvement
- Writing efficient and effective fuzz-tests will help benefit security and performance

- Nighthawk continued development & maturity
 - Continuous performance testing using public clouds
 - Visualization flow improvements
 - Promote as broader HTTP(2) measurement infrastructure
 - Contributors welcome: <https://github.com/envoyproxy/nighthawk>
- Fuzz testing for performance?
 - Semi-aligned with fuzz-testing for security bugs
 - Value increases as we make faster fuzzers
- **“Fast is Better than Slow!”**

Thank You